



# Explainable Federated Learning for Healthcare Time-Series Forecasting with Personalized Drift Adaptation

**Team Number: 28**

| S.No | Register Number  | Student Name           |
|------|------------------|------------------------|
| 1    | CB.SC.U4CSE23363 | HASWITHA K             |
| 2    | CB.SC.U4CSE23529 | HASINI REDDY M         |
| 3    | CB.SC.U4CSE23547 | SHEELA AKSHAR SAKHI    |
| 4    | CB.SC.U4CSE23761 | KOUSIK SARMA LAKKARAJU |
| 5    | CB.SC.U4CSE23753 | V. CHAKRAVARTHY        |

Guide: **DR. G R RAMYA, DR. VANDHANA S**

Department of CSE

# Previous Review Comments & Action Taken

## Reviewer Comments (Review 2)

- **Comment 1: Start Implementation**  
The implementation phase of the proposed federated framework must be initiated.
- **Comment 2: Working Prototype & Dashboard**  
Build a functional clinician dashboard and working prototype to demonstrate live clinical decision support.

## Action Taken & Implementation Status

- **Implemented Federated Pipelines:**
  - Built decentralized PyTorch algorithms (FedAvg, FedProx, Ditto).
  - Implemented local CUSUM drift detection and Client-Side Selective Personalization.
- **Built Production-Grade CDSS Workstation:**
  - **Data Warehouse:** Local SQLite clinical database.
  - **Backend API:** Uvicorn/FastAPI server running PyTorch checkpoints.
  - **Interactive UI:** Vite React medical dashboard mapping vitals charts, risk dials, and explainable attention timelines.

# Clinical Deterioration & Phase II Completed Goals

## Clinical Context:

- **High-Sparsity Time-Series:** Continuous ICU vital monitoring logs (HR, SBP, Temp, SpO<sub>2</sub>) combined with sparse blood panel entries.
- **Privacy constraints:** HIPAA / DPDPA (Section 8) regulations mandate that private patient records cannot leave local devices.
- **Concept Drifts:** Clinical demographics vary heavily across regional ICU nodes.

## SDG Alignment:

- **Primary (SDG 3):** Good Health & Well-being by detecting patient sepsis crashes early.
- **Secondary (SDG 9 & 16):** Secure and scalable digital clinical infrastructure.

## Phase II Completed Milestones

- **1. Data Engineering:** Cleansed, Imputed, Standardized PhysioNet 2019 logs into active PyTorch sequence arrays.
- **2. Federated Baselines:** Implemented *FedAvg* and *FedProx* securely over PyTorch.
- **3. Personalized FPDAF Model:** Trained temporal self-attentions with localized personalization.
- **4. Clinical Workstation App:** Developed a dynamic FastAPI + SQLite + React Vite ICU dashboard overlay.

# Clinical Profile of ICU Sepsis

## What Sepsis is

Sepsis is a life-threatening medical emergency caused by the body's extreme, systemic immune response to an infection, which ends up damaging its own tissues, blood vessels, and vital organs.

## Why it is Dangerous

If not diagnosed immediately, it rapidly progresses to Septic Shock, leading to multi-organ failure, a sudden drop in blood pressure, and high mortality rates.

## The Clinical Time-Window Challenge

Sepsis is notoriously difficult to detect early because its initial symptoms resemble general infections. However, every single hour of delayed treatment increases a patient's risk of death by approximately 8%.

## Role of Our Project

By continuously analyzing ICU vital signals, our FPDFAF early-warning system acts as an early-warning radar system, alerting doctors hours before a patient clinically deteriorates, giving clinicians a vital head-start to administer antibiotics and fluids.

# Sepsis Burden & Impact in India

## National Mortality & Case Load

Sepsis represents a critical public health crisis in India:

- **2.9 to 3.0 Million Deaths:** Estimated annually across India due to Sepsis.
- **11.0 to 11.3 Million Cases:** Recorded annually nationwide.
- **Proportional Mortality:** Sepsis accounts for roughly **30% (3 in 10)** of all deaths in India.

Source: [The George Institute for Global Health, 2020](#)

## Indian ICU Mortality Metrics

Within clinical critical care environments:

- **34% to 50%+ Mortality Rate:** Severe sepsis patients in Indian ICUs face high clinical mortality.
- **Antimicrobial Resistance (AMR):** High AMR rates complicate treatment, making hours-early detection crucial.

Source: [Indian Society of Critical Care Medicine \(ISCCM\)](#)

## Clinical Significance & Research Objective

Given that severe sepsis mortality in Indian ICUs exceeds 34%, implementing a localized, real-time early warning system is the most critical and clinically viable strategy to save patient lives.

# Proposed Framework: FPDAF

## WHAT is FPDAF?

A privacy-preserving Clinical Decision Support System (CDSS) for Sepsis early warning. Enables collaborative, HIPAA-compliant training across decentralized hospital ICU nodes without centralizing raw records.

## HOW does it work?

- **Sequence Modeling:** Extracts temporal clinical features using a shared Bidirectional LSTM.
- **Ditto Personalization:** Fine-tunes localized classifier heads ( $v_k$ ) to fit hospital demographics.
- **Online CUSUM Trigger:** Audits prediction loss residuals locally to detect population drift.

# Core Scientific & Technical Novelties

## 1. Multi-Head Attention Saliency (Clinical Explainability)

Integrates a 4-head query-key temporal self-attention mechanism to output dynamic sequence saliency weights, mapping exactly which hourly vitals contributed to Sepsis alarms at the bedside.

## 2. Online CUSUM Residual Neural Drift Monitor


Tracks local validation loss residuals round-by-round ( $S_r > 3.0$ ). Concept drift is detected entirely client-side, avoiding constant full-parameter aggregation cycles.

## 3. Client-Side Selective Personalization (CSSP)

Freezes the shared BiLSTM/attention layers upon drift triggers to train local classifier heads exclusively.

- Bypasses global parameter uploads, saving **38% in communication bandwidth** (1.52 GB vs 2.48 GB).
- Cuts client-side adaptation training time to **6 minutes** (down from 25 minutes).

# Unified System Design & Neural Formulations



**Figure 1:** FPDFAF Split-Weight Model Architecture.

## Decoupled bi-level Optimization Model

$$\hat{y}_t^k = h_{\phi^k} \left( \sum_{i=1}^t \alpha_i \cdot g_{\theta}(X_i^k) \right)$$

- $g_{\theta}$ : Shared global LSTM backbone ( $\theta$ ).
- $h_{\phi^k}$ : Local private classifier head ( $\phi^k$ ).

## Temporal Multi-Head Self-Attention

$$e_i = v_a^T \tanh(W_s s_i + b_a)$$

$$\alpha_i = \frac{\exp(e_i)}{\sum_{j=1}^t \exp(e_j)}$$

Attributes risk probabilities to specific ICU sequence hours.



# Proactive Drift Adaptation: CUSUM Neural Trigger

**Concept Drift Challenge:** A static global model degrades when client ICU patient distributions shift over time. Standard Personalized FL (e.g. Ditto) constantly aggregates and fine-tunes all nodes, resulting in immense network overhead.

## The CUSUM Residual Monitor Trigger:

- We monitor cumulative validation loss residuals ( $e_r$ ) over rounds:

$$e_r = \mathcal{L}_{val,r} - \mu_0$$

$$S_r = \max(0, S_{r-1} + e_r - \kappa)$$

- If  $S_r > H$  (threshold  $H = 3.0$ ), we declare institutional drift.
- Triggers **Client-Side Selective Personalization (CSSP)**: freezes global backbone layers  $\theta$  and adapts classification heads  $\phi^k$  locally.

## Bandwidth Savings Benefits

- Aggregation is bypassed during stable performance rounds.
- Communication bandwidth is saved by 38% compared to standard personalized systems (1.52 GB vs 2.48 GB).
- Local personalization head fine-tuning takes only 6 minutes (down from 25 minutes).



# Five-Way Model Benchmarking Results

We evaluated our framework against centralized baselines and standard federated learning algorithms:

| Model Configuration     | Accuracy      | Precision    | Recall        | F1-Score      | AUROC         | Comm. Bandwidth                 |
|-------------------------|---------------|--------------|---------------|---------------|---------------|---------------------------------|
| Centralized Baseline    | 75.09%        | 7.14%        | 72.17%        | 0.1299        | 0.8058        | 0 MB (N/A)                      |
| FedAvg Baseline         | 75.94%        | 6.94%        | 67.19%        | 0.1259        | 0.7844        | 1.24 GB                         |
| FedProx Baseline        | 78.85%        | 8.12%        | 60.64%        | 0.1432        | 0.7933        | 1.24 GB                         |
| Ditto (Personalized)    | 82.45%        | 8.98%        | 63.58%        | 0.1574        | 0.7989        | 2.48 GB                         |
| <b>FPDAF (Proposed)</b> | <b>86.51%</b> | <b>9.81%</b> | <b>65.33%</b> | <b>0.1706</b> | <b>0.8214</b> | <b>1.52 GB (CSSP saved 38%)</b> |

## Scientific Findings

- **Privacy & Performance:** FPDAF preserves HIPAA compliance while achieving a test AUROC of **\*\*0.8214\*\*** (outperforming centralized baseline).
- **Decentralized Superiority:** FPDAF outperforms all standard federated baselines (Accuracy **\*\*+10.57%\*\***, F1 **\*\*+4.47%\*\*** vs. FedAvg).

# Ablation Study Verification

We verified the contributions of CUSUM triggers, temporal attention weights, and localized personalization by running ablation tests:

| Ablation Configuration       | Accuracy      | Precision    | Recall        | F1-Score      | AUROC         |
|------------------------------|---------------|--------------|---------------|---------------|---------------|
| FPDAF w/o CUSUM Trigger      | 83.51%        | 8.51%        | 55.33%        | 0.1475        | 0.7603        |
| FPDAF w/o Temporal Attention | 80.01%        | 8.13%        | 52.81%        | 0.1412        | 0.7878        |
| FPDAF w/o Personalization    | 78.87%        | 8.05%        | 63.58%        | 0.1430        | 0.7887        |
| <b>Full FPDAF (Proposed)</b> | <b>86.51%</b> | <b>9.81%</b> | <b>65.33%</b> | <b>0.1706</b> | <b>0.8214</b> |

- **No Attention:** Stripping temporal attention query projections leads to poor clinical explainability and limits optimization capacity.
- **No CUSUM:** Disabling residual control limits leaves client nodes blind to data distribution drifts, leading to performance deterioration under heterogeneous ICU demographics.

# Clinical Workstation Software Architecture

**Database-Driven Design:** Instead of placeholders, we built a fully functional clinical warehouse:


- **SQLite Database:** Houses tables for patients, decoded 24h vitals, predictions, self-attention, and drift history.
- **FastAPI REST Endpoints:** Feeds data securely to the client UI.

## Custom User Portals:

- **Doctor Portal:** Focuses on patient-care operations (searchable bed lists, hourly vital monitors, attention graphs, and risk prediction alerts).
- **Admin Portal:** Focuses on machine learning telemetry (round parameter aggregations, CUSUM drift charts, and baseline comparisons).

## Startup Script

Coded a clean 'run\_app.sh' script to boot both FastAPI and React Vite concurrently, handling terminations on Ctrl+C.



## Doctor's Bedside Diagnostic View (Live Patient Care)

- **Timeline Telemetry:** Plots hourly ICU vital segments (HR, SBP, Temp, Resp, SpO<sub>2</sub>) using Recharts.
- **Live Predictions Dials:** Overlays risk probability scores for all federated baselines.
- **Attention Saliency Maps:** Attributes risk alerts to specific sequence hours (e.g. onset hours 16–22) for bedside interpretability.

## Admin's Federated Telemetry View (System Operations)

- **Drift Controls:** Plots CUSUM curves, tracking drift accumulation limits ( $S_r > 3.0$ ) and CSSP backbone freezes.
- **Round Tracker:** Animates client-server parameter upload/download loops using Framer Motion.
- **Design System:** Tailwind CSS v4 class-based toggling supporting light/dark medical navy workstation styles.

# Clinician Dashboard Interface (Doctor Portal)




Figure 2: Doctor's Bedside Diagnostic Workspace.

## Doctor Portal Core Modules

- **Searchable Bed Directory:** Displays patient demographics, ward bed mappings, classifications, and classifier confidence scores.
- **Hourly Vital Telemetry:** Renders 24-hour ICU vital timelines (HR, SBP, Temp, Resp, SpO<sub>2</sub>) using clinical ranges.
- **XAI Saliency Charts:** Plots Multi-Head Attention weights to interpret Sepsis alerts.
- **Risk Predictor Dials:** Compares live model forecasts across FedAvg, FedProx, Ditto, and FPDF models.

# Clinician Dashboard Interface (Admin Portal)




Figure 3: Admin's Machine Learning Telemetry Workspace.

## Admin Portal Core Modules

- **Federated Loop Animation:** Animates localized model weight uploads and global server distributions using Framer Motion.
- **CUSUM Drift Telemetry:** Tracks real-time concept drift curves ( $S_r > 3.0$ ) and selective personalization triggers.
- **LaTeX Equations & Ablations:** Details mathematical optimization equations and five-way benchmark metrics.

# Summary of Completed Phase II Works

- **Data Preprocessing Completed:** Decoded sparse vitals logs to construct clean sequence windows.
- **PyTorch Implementations Completed:** Trained standard baselines alongside our proposed FPDAF model.
- **Database Schema Integrated:** SQLite warehouse populated with real patients, hourly vitals, and predictions.
- **Clinical WORKSTATION Finished:** Production-quality React Vite interface with tailored user views.
- **Deliverables Pushed to Git:** Full Phase II codebase pushed to main repository.

## Roadmap Compliance

Our completed milestones match the project objectives, delivering a privacy-preserving and explainable Sepsis prediction pipeline.



# References I



E. Churaev, I. Kisel, and A. Savchenko, "A standalone software for real-time facial analysis in online conferences and e-lessons," *Software Impacts*, 2023, mTCNN face detection; EfficientNet-B0 for FER; Apache TVM optimization.



R. Jothimurugesan, P. S. Hiremath, and S. Natarajan, "Concept drift detection and adaptation for federated and continual learning," *Multimedia Tools and Applications*, 2022, cDA-FedAvg with CUSUM drift detection.



A. Arivazhagan, V. Aggarwal, A. Bajpai, and S. Khudanpur, "Federated learning with personalization layers," in *Proceedings of AISTATS*. Elsevier, 2020, fedPer with shared base layers and local personalized layers.



H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, "Communication-efficient learning of deep networks from decentralized data," in *Proceedings of AISTATS*. PMLR, 2017, fedAvg using local SGD and global model averaging.



Y. Liu, X. Chen, Y. Liu, and Z. Luo, "Federated learning with dual-end gradient correction and proxy-free self-distillation," *IEEE Access*, 2022, dual-end gradient correction; stable federated learning under non-IID data.



X. Yang, J. Jung, J. Lu, K.-W. Lim, and L. Linguaglossa, "Mvfl – multivariate vertical federated learning for time-series forecasting," in *IFIP CNSM Conference Proceedings*, 2025, mVFL framework; vertical federated learning.



R. Kalakoti, S. Nomm, and H. Bahsi, "Fedxai: Federated learning of explainable ai for intrusion detection," *Computer Networks*, 2025, fedXAI with SHAP and LIME-based explanations.



M. Nosedà, A. D. Luca, L. V. Briel, and N. Lacour, "Federated learning for financial forecasting," ETH Zurich Technical Report / arXiv, 2025, federated LSTM models with client-side fine-tuning.



A. Karamanlioğlu, B. Demirel, O. Tural, O. T. Doğan, and F. N. Alpaslan, "Privacy-preserving clinical decision support for emergency triage using llms," *MDPI*, 2025, federated learning with LLMs; Scopus Indexed.

# References II



R. Wang, L. Bai, R. Rayhana, and Z. Liu, "Personalized federated learning for time-series forecasting," *Elsevier / ACM*, 2024–2025, shared global model with client-specific heads.



R. Abdel-Sater and A. B. Hamza, "A federated large language model for long-term time series forecasting," arXiv, 2024, federated learning with LLaMA; parameter-efficient fine-tuning.



M. Maher, O. F. Oun, M. S. Mesmeh, and R. ElShawi, "Fedforecaster: An automated federated learning approach for time-series forecasting," in *EDBT Conference*, 2025, autoML-based federated learning; Bayesian optimization.



S. Gupta and V. Torra, "Privacy-enhancing federated time-series forecasting: A microaggregation-based approach," in *SECRYPT*, 2025, microaggregation with k-anonymity.



W. Yuan, G. Ye, X. Zhao, Q. V. H. Nguyen, Y. Cao, and H. Yin, "Tackling data heterogeneity in federated time series forecasting," arXiv, 2024, fed-TREND framework; synthetic data generation.



D. Tang, N. Yang, Y. Li, Z. Zhu, Z. Jin, and D. Yuan, "Optimal look-back horizon for time series forecasting in federated learning," AAAI / arXiv, 2025, adaptive look-back horizon selection.



Y. Zhang, J. Wang, H. Liu, and X. Chen, "Federated multi-task learning for time series forecasting," *IEEE Transactions on Neural Networks and Learning Systems*, 2022, federated multi-task learning; shared global representation.



G. L. Rocca, A. S. Podda, R. Martoglia, and M. Botta, "Interpretable multivariate time-series forecasting with attention mechanisms," in *NeurIPS*, 2021, attention-based encoder-decoder; temporal and variable-level attention.



G. Amato, C. Gennaro, A. Matteucci, and L. Serafini, "Federated learning under heterogeneous time series data," in *ACM SIGKDD*, 2023, client clustering; similarity-based grouping.

# References III



G. L. Rocca, A. S. Podda, R. Martoglia, and M. Botta, "Explainable ai for multivariate time series forecasting," *Information Sciences*, 2022, SHAP and LIME explanations; model-agnostic interpretability.



M. Bernardi, P. Giordani, and F. Ravazzolo, "Personalized federated learning with adaptive model fusion," in *AAAI*, 2023, adaptive model fusion; client-level weighting.



C. Düsing and P. Cimiano, "Early sepsis onset prediction through federated learning," arXiv, 2025, sliding-window strategy with Attention-LSTM; FedAvg aggregation.



A. Qayyum, K. Ahmad, M. A. Ahsan, A. Al-Fuqaha, and J. Qadir, "Fedsepsis: A federated multi-modal deep learning-based application," *MDPI*, 2023, edge computing with CNN-LSTM; GAN-based data augmentation.



J. Lei, J. Zhai, Y. Zhang, J. Qi, and C. Sun, "Supervised ml models for predicting sepsis and decision support," *JMIR Medical Informatics*, 2025, stacking ensemble (XGBoost); SHAP-based feature importance.



W. Pan, Z. Xu, S. Rajendran, and F. Wang, "Multiprog: Privacy-preserving fl framework," *PMC*, 2025, multi-channel federated learning; feature calibration.



A. Maurya, R. Haripriya, M. Pandey, J. Choudhary, and D. P. Singh, "Federated learning for severity classification at the edge," *IEEE Access*, 2025, secure multi-party computation; encrypted FedAvg.

Thank You Questions?