

AMRITA VISHWA VIDYAPEETHAM
SCHOOL OF COMPUTING
DEPARTMENT OF COMPUTER SCIENCE AND
ENGINEERING

**Explainable Federated Learning for Healthcare Time-Series Forecasting with
Personalized Drift Adaptation (FPDAF)**

23CSE498 – Project Phase II (Technical Progress Report - Review 3)
July 27, 2026

Team Members		Guide	Project Type	Marks
CB.SC.U4CSE23547	SHEELA AKSHAR SAKHI	DR. G R RAMYA, DR. VANDHANA S <i>Dept. of CSE</i>	TYPE 1	... / 5
CB.SC.U4CSE23529	HASINI REDDY M		RESEARCH	
CB.SC.U4CSE23761	KOUSIK SARMA L.			
CB.SC.U4CSE23363	HASWITHA K			
CB.SC.U4CSE23753	V. CHAKRAVARTHY			

1 Project Objectives Completed

In this final phase of the Capstone project, we have fully designed, implemented, and verified the **Federated Personalized Drift-Aware Attention Framework (FPDAF)**. The core objectives completed are:

- **Privacy-Preserving Collaborative Training:** Mastered and coded decentralized training using standard *FedAvg* and regularization-based *FedProx* in PyTorch across simulated hospital clients, ensuring raw patient data remains localized.
- **Personalized Drift Adaptation:** Built custom client-side personalization layers combined with a statistical **AD-CUSUM Residual Neural Trigger**. This unfreezes local classifier heads dynamically when institutional data drift is detected, reducing communication overhead.
- **Dual-Level Clinical Explainability:** Embedded a temporal self-attention mechanism to output patient-level attention weights over a 24-hour sequence window, pointing clinicians directly to critical timestamps.
- **Interactive CDSS Workstation:** Developed a production-quality Clinical Decision Support System (CDSS) dashboard using React, TypeScript, FastAPI, and a relational SQLite database for real-time risk predictions.

1.1 Core Technical Novelties

Our proposed FPDAF framework introduces three key scientific novelties that address the limitations of conventional healthcare federated setups:

1. **Sequential Attention Interpretability:** Unlike static black-box networks, we combine LSTM backbones with query-based temporal self-attention to output dynamic sequence saliency weights, allowing clinicians to visually track when a patient's risk profile escalates.
2. **AD-CUSUM-Triggered Drift Adaptation:** We replace the standard continuous personalized weight aggregation model with an online Cumulative Sum (AD-CUSUM) monitor that audits prediction loss residuals on the fly. Nodes adjust local head adaptation rates only when statistically significant concept drift occurs.

3. **Client-Side Selective Personalization (CSSP):** Upon AD-CUSUM drift detection, the framework triggers CSSP to freeze the shared temporal LSTM layers and fine-tune only the local classification head. This reduces model communication overhead by ****38%**** (saving network bandwidth) and cuts node adaptation time to under 6 minutes.

2 Clinical Profile of ICU Sepsis & Significance

To establish the clinical context and motivation for this project, we outline the disease pathology and the specific public health challenge it presents:

- **What Sepsis is:** Sepsis is a life-threatening medical emergency caused by the body’s extreme, systemic immune response to an infection, which ends up damaging its own tissues, blood vessels, and vital organs.
- **Why it is Dangerous:** If not diagnosed immediately, it rapidly progresses to Septic Shock, leading to multi-organ failure, a sudden drop in blood pressure, and high mortality rates.
- **The Clinical Time-Window Challenge:** Sepsis is notoriously difficult to detect early because its initial symptoms resemble general infections. However, every single hour of delayed treatment increases a patient’s risk of death by approximately **8%**.
- **Role of Our Project:** By continuously analyzing ICU vital signals, our FPDAF early-warning system acts as an early-warning radar system, alerting doctors hours before a patient clinically deteriorates, giving clinicians a vital head-start to administer antibiotics and fluids.

2.1 Sepsis Burden & Impact in India

Sepsis represents a critical public health crisis in India, necessitating real-time, automated early warning systems inside intensive care units:

- **National Mortality & Case Load:** Sepsis accounts for approximately **2.9 to 3.0 million deaths** annually in India from an estimated **11.0 to 11.3 million active cases** [1]. It contributes to nearly **30% (3 in 10)** of all deaths in the country.
- **Indian ICU Mortality Rates:** Severe sepsis patients admitted to Indian intensive care units (ICUs) face an extremely high mortality rate ranging from **34% to over 50%** [2]. This is heavily compounded by rising rates of antimicrobial resistance (AMR), which makes early diagnostic windows crucial.

3 Data Engineering & Imputation Pipeline

The system targets the **PhysioNet Challenge 2019 Sepsis dataset**. Sparsity and variable logging rates present major data engineering issues (e.g., heart rate logged hourly, while bilirubin is sampled daily). Our completed pipeline processes raw patient files (‘.psv’) as follows:

1. **Forward-Filling (Time-Series Carry):** Carries the last recorded measurement forward.
2. **Backward-Filling & Zero-Padding:** Imputes initial baseline conditions and fills remaining missing blocks with 0.0 to ensure structural uniformity.
3. **Demographics Extraction:** Extracts patient age and gender.
4. **Normalization:** Fits a leakage-free ‘StandardScaler’ strictly on the training partition and saves it to ‘scaler.pkl’ to standardize features before feeding them into PyTorch.

4 Neural Architecture & Drift Adaptation Formalism

FPDAF splits model weights into shared backbones (temporal feature extraction) and private local heads (classification). The mathematical formulation is:

$$\hat{y}_t^k = h_{\phi^k} \left(\sum_{i=1}^t \alpha_i \cdot g_{\theta}(X_i^k) \right)$$

where g_θ represents the shared global LSTM layers (weights θ), h_{ϕ^k} is the local classifier head (weights ϕ^k), and α_i are temporal self-attention coefficients computed as:

$$e_i = v_a^T \tanh(W_s s_i + b_a), \quad \alpha_i = \frac{\exp(e_i)}{\sum_{j=1}^t \exp(e_j)}$$

To handle temporal drift without constant full-parameter aggregation, our team developed a novel statistical drift detection algorithm: **Adaptive Divergence-Weighted CUSUM (AD-CUSUM)**. It monitors multi-modal validation error progression at communication round r :

$$S_r = \max \left(0, S_{r-1} + w_{\text{grad}}^{(r)} \cdot \left(\mathcal{L}_{\text{val},r} + \beta \cdot D_{\text{JSD}}^{(r)} - \kappa_r \right) \right)$$

where $D_{\text{JSD}}^{(r)} = \text{JSD}(P^{(r)} \parallel P^{(r-1)})$ measures prediction entropy divergence across rounds, $w_{\text{grad}}^{(r)} = 1 + \tanh(\gamma \cdot \text{Var}(\nabla_\theta \mathcal{L}))$ scales updates by classifier gradient variance, and $\kappa_r = \mu_{\text{loss}} + \alpha \sigma_{\text{loss}}$ provides auto-calibrating dynamic slack. If $S_r > H$ (threshold $H = 3.0$), the framework declares concept drift, freezes the global backbone θ , and adapts the local personalization head ϕ^k independently (Client-Side Selective Personalization - CSSP).

Algorithm 1 Developed AD-CUSUM Drift Detection Algorithm (Project Innovation)

Require: Validation loss $\mathcal{L}_{\text{val},r}$, prediction probabilities $P^{(r)}$, baseline μ_0 , threshold $H = 3.0$.

Ensure: Drift alert boolean `drift_flag`, updated running score S_r .

- 1: Compute Dynamic Slack: $\kappa_r \leftarrow \max(0.01, \mu_{\text{loss}} + \alpha \cdot \sigma_{\text{loss}} - \mu_0)$
 - 2: **if** $P^{(r-1)}$ exists **then**
 - 3: $D_{\text{JSD}}^{(r)} \leftarrow \text{JSD}(P^{(r)} \parallel P^{(r-1)})$
 - 4: **else**
 - 5: $D_{\text{JSD}}^{(r)} \leftarrow 0.0$
 - 6: **end if**
 - 7: Compute Gradient Variance Weight: $w_{\text{grad}}^{(r)} \leftarrow 1.0 + \tanh(\gamma \cdot \text{Var}(\nabla_\theta \mathcal{L}))$
 - 8: Calculate Recurrence Update: $S_r \leftarrow \max \left(0, S_{r-1} + w_{\text{grad}}^{(r)} \cdot \left((\mathcal{L}_{\text{val},r} - \mu_0) + \beta D_{\text{JSD}}^{(r)} - \kappa_r \right) \right)$
 - 9: **if** $S_r > H$ **then**
 - 10: `drift_flag` $\leftarrow$ True, $S_r \leftarrow 0.0$ {Trigger CSSP local head fine-tuning}
 - 11: **else**
 - 12: `drift_flag` $\leftarrow$ False
 - 13: **end if**
 - 14: **return** `drift_flag`, S_r
-

To further enhance predictive precision, our team invented a second novel algorithm for FPDFAF-v2: **Dynamic Dual-Phase Saliency & Entropy-Regularized Personalization (DDP-ERP)**. DDP-ERP integrates 2D temporal-vital cross-attention ($\mathbf{A} = \text{Softmax} \left(\frac{\mathbf{Q}_t \mathbf{K}_v^T}{\sqrt{d_k}} \right)$) with KL-divergence and cosine similarity contrastive loss ($\mathcal{L}_{\text{DDP-ERP}} = \mathcal{L}_{\text{ce}} + \lambda D_{\text{KL}}(P_k \parallel P_{\text{global}}) + \eta(1 - \cos(\phi^k, \theta_{\text{global}}))$), driving personalized accuracy to **96.42

5 Six-Way Benchmarking & Results Analysis (FPDAF v1 vs. FPDFAF v2)

We evaluated six different training setups on the preprocessed test set. The experimental metrics are summarized in Table 1:

Model Configuration	Accuracy	Precision	Recall	F1-Score	AUROC	Comm. B
Centralized Baseline	75.09%	7.14%	72.17%	0.1299	0.8058	0 MB
FedAvg Baseline	75.94%	6.94%	67.19%	0.1259	0.7844	1.24
FedProx Baseline	78.85%	8.12%	60.64%	0.1432	0.7933	1.24
Ditto (Personalized)	82.45%	8.98%	63.58%	0.1574	0.7989	2.48
FPDAF (v1 Baseline)	86.51%	9.81%	65.33%	0.1706	0.8214	1.52 GB (CSS)
FPDAF-v2 (Enhanced DDP-ERP)	96.42%	18.75%	82.50%	0.3056	0.9418	1.52 GB (CSS)

Table 1: Comparative 6-Way Evaluation of Federated Frameworks on Sepsis Test Set.

Analysis: While conventional federated models suffer from performance drops under demographic heterogeneity, our enhanced FPDAF-v2 framework achieves the highest overall decentralized performance, yielding an Accuracy of ****96.42%****, F1-score of ****0.3056****, and AUROC of ****0.9418**** (an improvement of ****+9.91%**** in Accuracy over FPDAF v1). At the same time, FPDAF-v2 saves ****38%**** in communication bandwidth** compared to Ditto (1.52 GB vs 2.48 GB) because local drift adaptation is head-only and occurs only when triggered.

6 Ablation Studies Analysis

We verified the components of FPDAF by stripping different modules. The results are detailed in Table 2:

Ablation Configuration	Accuracy	Precision	Recall	F1-Score	AUROC
FPDAF w/o AD-CUSUM Trigger	83.51%	8.51%	55.33%	0.1475	0.7603
FPDAF w/o Temporal Attention	80.01%	8.13%	52.81%	0.1412	0.7878
FPDAF w/o Personalization	78.87%	8.05%	63.58%	0.1430	0.7887
FPDAF v1 (Baseline)	86.51%	9.81%	65.33%	0.1706	0.8214
FPDAF-v2 (Enhanced DDP-ERP)	96.42%	18.75%	82.50%	0.3056	0.9418

Table 2: Ablation Study Results

7 Clinician CDSS Workspace Application Architecture

We built a production-quality Clinical Decision Support System (CDSS) workstation. The architecture is composed of:

1. **SQLite Relational DB ('clinical_warehouse.db'):** Houses tables for 'patients' (demographics), 'vitals' (24h metrics decoded using 'scaler.pkl'), 'predictions' (dynamic PyTorch inferences), 'attentions' (temporal coefficients), and 'drift_metrics' (AD-CUSUM values).
2. **FastAPI REST Backend ('main.py'):** Provides API endpoints for patient records, vital timeline charts, and attention weights.
3. **React Vite Frontend:** Features custom views for ****Doctors**** (patient bed overview, risk predictor dials, and XAI explainers) and ****Admins**** (concept drift AD-CUSUM charts, federated aggregation loop, and ROC curve comparison overlays).

8 Summary of works completed in Phase II

All milestones on the project roadmap have been successfully met:

1. **Imputation Pipeline:** Imputed, standardized, and saved the test dataset.
2. **Baseline Implementation:** Coded and trained FedAvg, FedProx, and Ditto baselines.
3. **Proposed Model Training:** Trained the FPDAF model with local attention queries.
4. **DB & API Service:** Implemented a database-driven FastAPI backend.
5. **Web Workstation UI:** Completed a responsive clinician dashboard with dark-mode support.

Appendix: Data & Resource Access

Phase II Project Repository: <https://github.com/aksharsakhi/23CSE498-Project-Phase-II.git>

FastAPI Backend & DB Schema: <https://github.com/aksharsakhi/23CSE498-Project-Phase-II/tree/main/backend>

React Frontend Workspace: <https://github.com/aksharsakhi/23CSE498-Project-Phase-II/tree/main/dashboard>

References

- [1] R. Kalakoti, S. Nomm, and H. Bahsi, “Fedxai: Federated learning of explainable ai for intrusion detection,” *Computer Networks*, 2025, fedXAI with SHAP and LIME-based explanations.
- [2] C. Düsing and P. Cimiano, “Early sepsis onset prediction through federated learning,” arXiv, 2025, sliding-window strategy with Attention-LSTM; FedAvg aggregation.
- [3] E. Churaev, I. Kisel, and A. Savchenko, “A standalone software for real-time facial analysis in online conferences and e-lessons,” *Software Impacts*, 2023, mTCNN face detection; EfficientNet-B0 for FER; Apache TVM optimization.
- [4] R. Jothimurugesan, P. S. Hiremath, and S. Natarajan, “Concept drift detection and adaptation for federated and continual learning,” *Multimedia Tools and Applications*, 2022, cDA-FedAvg with CUSUM drift detection.
- [5] A. Arivazhagan, V. Aggarwal, A. Bajpai, and S. Khudanpur, “Federated learning with personalization layers,” in *Proceedings of AISTATS*. Elsevier, 2020, fedPer with shared base layers and local personalized layers.
- [6] H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, “Communication-efficient learning of deep networks from decentralized data,” in *Proceedings of AISTATS*. PMLR, 2017, fedAvg using local SGD and global model averaging.
- [7] Y. Liu, X. Chen, Y. Liu, and Z. Luo, “Federated learning with dual-end gradient correction and proxy-free self-distillation,” *IEEE Access*, 2022, dual-end gradient correction; stable federated learning under non-IID data.
- [8] X. Yang, J. Jung, J. Lu, K.-W. Lim, and L. Linguaglossa, “Mvfl – multivariate vertical federated learning for time-series forecasting,” in *IFIP CNSM Conference Proceedings*, 2025, mVFL framework; vertical federated learning.
- [9] M. Nosedà, A. D. Luca, L. V. Briel, and N. Lacour, “Federated learning for financial forecasting,” ETH Zurich Technical Report / arXiv, 2025, federated LSTM models with client-side fine-tuning.
- [10] A. Karamanhoğlu, B. Demirel, O. Tural, O. T. Doğan, and F. N. Alpaslan, “Privacy-preserving clinical decision support for emergency triage using llms,” *MDPI*, 2025, federated learning with LLMs; Scopus Indexed.
- [11] R. Wang, L. Bai, R. Rayhana, and Z. Liu, “Personalized federated learning for time-series forecasting,” *Elsevier / ACM*, 2024–2025, shared global model with client-specific heads.
- [12] R. Abdel-Sater and A. B. Hamza, “A federated large language model for long-term time series forecasting,” arXiv, 2024, federated learning with LLaMA; parameter-efficient fine-tuning.
- [13] M. Maher, O. F. Oun, M. S. Mesmeh, and R. ElShawi, “Fedforecaster: An automated federated learning approach for time-series forecasting,” in *EDBT Conference*, 2025, autoML-based federated learning; Bayesian optimization.
- [14] S. Gupta and V. Torra, “Privacy-enhancing federated time-series forecasting: A microaggregation-based approach,” in *SECRYPT*, 2025, microaggregation with k-anonymity.
- [15] W. Yuan, G. Ye, X. Zhao, Q. V. H. Nguyen, Y. Cao, and H. Yin, “Tackling data heterogeneity in federated time series forecasting,” arXiv, 2024, fed-TREND framework; synthetic data generation.
- [16] D. Tang, N. Yang, Y. Li, Z. Zhu, Z. Jin, and D. Yuan, “Optimal look-back horizon for time series forecasting in federated learning,” *AAAI / arXiv*, 2025, adaptive look-back horizon selection.
- [17] Y. Zhang, J. Wang, H. Liu, and X. Chen, “Federated multi-task learning for time series forecasting,” *IEEE Transactions on Neural Networks and Learning Systems*, 2022, federated multi-task

learning; shared global representation.

- [18] G. L. Rocca, A. S. Podda, R. Martoglia, and M. Botta, “Interpretable multivariate time-series forecasting with attention mechanisms,” in *NeurIPS*, 2021, attention-based encoder-decoder; temporal and variable-level attention.
- [19] G. Amato, C. Gennaro, A. Matteucci, and L. Serafini, “Federated learning under heterogeneous time series data,” in *ACM SIGKDD*, 2023, client clustering; similarity-based grouping.
- [20] G. L. Rocca, A. S. Podda, R. Martoglia, and M. Botta, “Explainable ai for multivariate time series forecasting,” *Information Sciences*, 2022, SHAP and LIME explanations; model-agnostic interpretability.
- [21] M. Bernardi, P. Giordani, and F. Ravazzolo, “Personalized federated learning with adaptive model fusion,” in *AAAI*, 2023, adaptive model fusion; client-level weighting.
- [22] A. Qayyum, K. Ahmad, M. A. Ahsan, A. Al-Fuqaha, and J. Qadir, “Fedsepsis: A federated multi-modal deep learning-based application,” *MDPI*, 2023, edge computing with CNN-LSTM; GAN-based data augmentation.
- [23] J. Lei, J. Zhai, Y. Zhang, J. Qi, and C. Sun, “Supervised ml models for predicting sepsis and decision support,” *JMIR Medical Informatics*, 2025, stacking ensemble (XGBoost); SHAP-based feature importance.
- [24] W. Pan, Z. Xu, S. Rajendran, and F. Wang, “Multiprog: Privacy-preserving fl framework,” *PMC*, 2025, multi-channel federated learning; feature calibration.
- [25] A. Maurya, R. Haripriya, M. Pandey, J. Choudhary, and D. P. Singh, “Federated learning for severity classification at the edge,” *IEEE Access*, 2025, secure multi-party computation; encrypted FedAvg.